

Windows Kernel Reference Count Vulnerabilities - Case Study

Mateusz "j00ru" Jurczyk



@ ZeroNights E.0x02
November 2012

```
PS C:\Users\j00ru> whoami
```

```
nt authority\system
```

- Microsoft Windows internals fanboy
- Also into reverse engineering and low-level software security
- Currently in Switzerland working at Google

Why this talk?

- Lost of stuff in a sandbox
 - Google Chrome, Adobe Reader, Apple Safari, pepper plugins, ...
 - Escapes are becoming valuable
- Also, escapes are super exciting!
 - <https://krebsonsecurity.com/2012/11/experts-warn-of-zero-day-exploit-for-adobe-reader/> (just recently)
 - ... really, is this so shocking?
- "New" old class of bugs in the Windows kernel
- Otherwise, a bunch of technically interesting bugs


Topics covered

- Reference counting philosophy and problems
- Case study
 - a. 1-day (NT Object Manager *PointerCount* weakness)
 - b. 0-day (generic device driver image use-after-free)
 - c. CVE-2010-2549 (win32k!NtUserCheckAccessForIntegrityLevel use-after-free)
 - d. CVE-2012-2527 (win32k!NtUserAttachThreadInput use-after-free)
 - e. CVE-2012-1867 (win32k!NtGdiAddFontResource use-after-free)
- Mitigations and lessons learned

Reference counting

Fundamentals

- From now on, considering ring-0 refcounting
- System state → graph
 - resources → nodes
 - dependencies (refs) → directed edges
 - lonely node → destroy
 - dynamic memory management = vulnerabilities



Fundamentals

- In the graph scenario, a vertex doesn't have to know who points at him
 - Just the **total number**

- Common expression in garbage collectors:

```
if (!pObject->Refcount) {  
    free(pObject);  
}
```

- Unsurprisingly, refcounting is usually implemented using **plain integers**

Fundamentals

- Typical code pattern

```
POBJECT pObject = TargetObject;  
PCLIENT pClient = ClientObject;
```

```
pObject->Refcount++;
```

```
pClient->InternalPtr = pObject;
```

```
/* Perform operations on pClient assuming  
initialized InternalPtr */
```

```
pClient->InternalPtr = NULL;
```

```
pObject->Refcount--;
```

pObject guaranteed to persist



Fundamentals

- Windows kernel primarily written in C
- Everything is (described by) a structure
- Lack of common interface to manage references
 - Implemented from scratch every single time when needed...
 - ... always in a different way

Examples?

```
kd> dt _OBJECT_HEADER
nt!_OBJECT_HEADER
+0x000 PointerCount      : Int8B
+0x008 HandleCount      : Int8B
+0x008 NextToFree       : Ptr64 Void
+0x010 Lock              : _EX_PUSH_LOCK
[...]
```

```
kd> dt _LDR_DATA_TABLE_ENTRY
nt!_LDR_DATA_TABLE_ENTRY
[...]
```

| | | |
|--------|-----------|---------------|
| +0x068 | Flags | : Uint4B |
| +0x06c | LoadCount | : Uint2B |
| +0x06e | TlsIndex | : Uint2B |
| +0x070 | HashLinks | : _LIST_ENTRY |

```
[...]
```

```
kd> dt tagQ
win32k!tagQ
+0x000 mlInput           : tagMLIST
[...]
```

| | | |
|--------|-----------------|-------------------|
| +0x070 | hwndDblClk | : Ptr64 HWND_ |
| +0x078 | ptDblClk | : tagPOINT |
| +0x080 | ptMouseMove | : tagPOINT |
| +0x088 | afKeyRecentDown | : [32] UChar |
| +0x0a8 | afKeyState | : [64] UChar |
| +0x0e8 | caret | : tagCARET |
| +0x130 | spcurCurrent | : Ptr64 tagCURSOR |
| +0x138 | iCursorLevel | : Int4B |
| +0x13c | QF_flags | : Uint4B |
| +0x140 | cThreads | : Uint2B |
| +0x142 | cLockCount | : Uint2B |

```
[...]
```

Reference counting: problems

Logical issues

- Crucial requirement: refcount **must** be adequate to number of references by pointer
- Obviously, two erroneous conditions
 - Refcount is inadequately **small**
 - Refcount is inadequately **large**
- Depending on the context, both may have serious implications

Overly small refcounts


- Two typical reasons
 - Reference-by-pointer without refcount incrementation
 - More decrementations in a *destroy* phase than incrementations performed before
- Foundation of modern user-mode vulnerability hunting (web browsers et al)
 - <http://zerodayinitiative.com/advisories/published/>
 - <http://blog.chromium.org/2012/06/tale-of-two-pwnies-part-2.html>
 - <https://www.google.pl/#q=metasploit+use-after-free>
 - ...

Overly small refcounts

- Typical outcome in ring-3

```
mov eax, dword ptr [ecx]
mov edx, dword ptr [eax+70h]
call edx
```

object vtable lookup + call



- Still use-after-free in ring-0, but not so trivial
 - almost no vtable calls in kernel
 - exploitation of each case is bug specific and usually requires a lot of work
 - kernel pools feng shui is far less developed and documented compared to userland
 - Tarjei Mandt has exploited a few, check his BH slides and white-paper

Overly large refcounts

- Expected result → resource is never freed
 - Memory leak
 - Potential DoS via memory exhaustion
 - Not very useful
- But refcounts are integers, remember?
 - Finite precision.
 - Integer arithmetic problems apply!
 - Yes, we can try to overflow
- This can become a typical "small refcount" problem
 - use-after-free again

Reference count leaks

- If we can trigger a leak for free, it's exploitable

```
while (1) {  
    TriggerRefCountLeak(pObject);  
}
```

- Unless the integer range is too large
 - uint16_t is not enough
 - uint32_t is (usually) not enough anymore
 - uint64_t is enough

Reference count leaks

- Or unless object pinning implemented
(`ntdll!LdrpUpdateLoadCount2`)

```
if (Entry->LoadCount != 0xffff) {  
    // Increment or decrement the refcount  
}
```



j00ru/vx
@j00ru

Interesting Windows behavior: once you load a DLL 65535 times via LoadLibrary, it will stay there forever (see LdrpLoadDll / LdrpUnloadDll)

← Reply 🗑 Delete ★ Favorite

Legitimately large refcounts

- Sometimes even those can be a problem
- We can bump up refcounts up to a specific value
- Depends on bound memory allocations

never happens

| Per-iteration byte limit | Reference counter size |
|------------------------------|------------------------|
| impossible | 64 bits |
| 0-2 bytes | 32 bits |
| 16,384 - 131,072 bytes | 16 bits |
| 4,194,304 - 33,554,432 bytes | 8 bits |

Perfect reference counting

Qualities

- **Implementation:** 32-bit or 64-bit (safe choice) integers.
- **Implementation:** sanity checking, e.g.
`refcount ≥ 0x80000000 ⇒ bail out`
- **Usage:** `reference# = dereference#`
 - Random idea: investigate system state at shutdown
- **Usage:** never use object outside of its reference block
- **Mitigation:** *reference typing*

Reference counting bugs: case study

NT Object Manager PointerCount weakness

- Manages common resources
 - files, security tokens, events, mutants, timers, ...
 - around 50 types in total (most very obscure)
- Provides means to (de)reference objects
 - Public kernel API functions
 - ObReferenceObject, ObReferenceObjectByHandle, ObReferenceObjectByHandleWithTag, ObReferenceObjectByPointer, ObReferenceObjectByPointerWithTag, ObReferenceObjectWithTag
 - ObDereferenceObject, ObDereferenceObjectDeferDelete, ObDereferenceObjectDeferDeleteWithTag, ObDereferenceObjectWithTag
 - Extensively used by the kernel itself and third-party drivers

NT Object Manager PointerCount weakness

Fundamentals

- Each object comprised of a header + body
 - Header common across all objects, body specific to type (e.g. ETHREAD, EPROCESS, ERESOURCE)

```
kd> dt _OBJECT_HEADER
win32k!_OBJECT_HEADER
    +0x000 PointerCount
    +0x004 HandleCount
    [...]
    +0x008 Type
    [...]
    +0x018 Body
```

: Int4B

: Int4B

: Ptr32 _OBJECT_TYPE

: _QUAD

native word-wide
reference counters

type specifier

type specific structure

NT Object Manager PointerCount weakness

Fundamentals

- Two reference counters
 - `PointerCount` - # of direct kernel-mode pointer references
 - `HandleCount` - # of indirect references via `HANDLE` (both ring-3 and ring-0)
- Object free condition `(PointerCount == 0) && (HandleCount == 0)`

NT Object Manager PointerCount weakness

- Security responsibility put on the caller
 - Allows arbitrary number of decrementations
 - Allows reference count integer overflows
- Excessive dereferences rather uncommon
 - CVE-2010-2549 is the only I can remember
- Reference leaks on the other hand...
 - can *theoretically* only lead to memory leak
 - who'd care?
 - *sometimes you just forget to close something*
 - much more popular (in third-parties, not Windows)

NT Object Manager PointerCount weakness

- Userland can't overflow HandleCount
 - At least 32GB required to store four billion descriptors.
 - HANDLE address space is four times smaller than a native word.
- But random drivers can overflow PointerCount
 - grep through %system32%\drivers?

```
< Binary file ./cpqdap01.sys matches  
< Binary file ./isapnp.sys matches  
< Binary file ./modem.sys matches  
< Binary file ./nwlnkipx.sys matches  
< Binary file ./pcmcia.sys matches  
< Binary file ./sdbus.sys matches  
< Binary file ./wmilib.sys matches
```



Import a *Reference*,
but no *Dereference*
symbol.

NT Object Manager PointerCount weakness

- Refcount leaks are as dangerous as double derefs (only on 32-bit platforms)
 - just take longer to exploit
- Had a chat with *Microsoft security*
- A few months later, Windows 8 ships with a fix:

```
[...]  
v8 = _InterlockedIncrement((signed __int32 *)v5);  
if ( (signed int)v8 <= 1 )  
    KeBugCheckEx(0x18u, 0, ObjectBase, 0x10u, v8);  
[...]
```

" The REFERENCE_BY_POINTER bug check has a value of 0x00000018. This indicates that the reference count of an object is illegal for the current state of the object. "

NT Object Manager PointerCount weakness

- Ken Johnson and Matt Miller covered this and other mitigations during their BH USA 2012 presentation
 - ["Exploit Mitigation Improvements in Windows 8"](#), check it out
- Mitigation only released for Windows 8
 - older platforms still affected
 - go and find your own unpaired *ObReferenceObject* invocations?

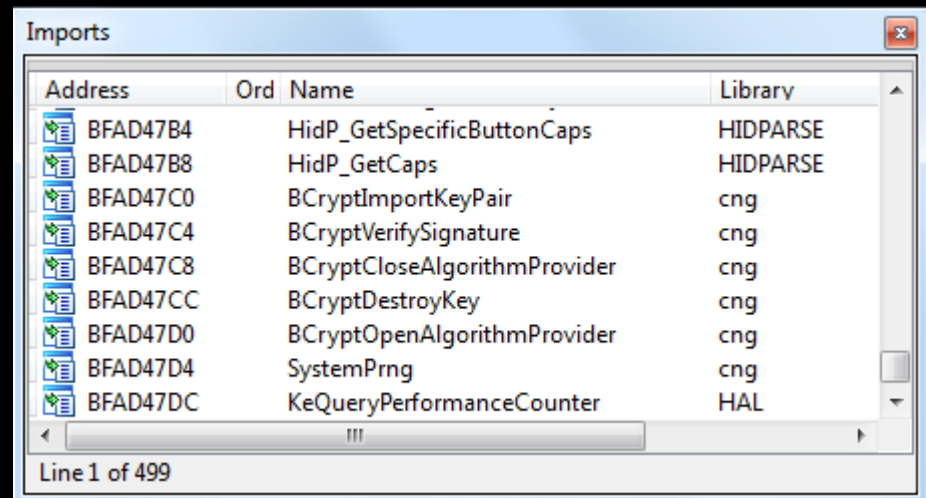
Device driver image use-after-free

- Many drivers loaded in Windows at any time

```
kd> lm
```

| start | end | module name | |
|----------|----------|-------------|---------------|
| 80ba0000 | 80ba8000 | kdcom | (deferred) |
| 8281f000 | 82c31000 | nt | (pdb symbols) |
| 82c31000 | 82c68000 | hal | (deferred) |
| 82e00000 | 82e25000 | CLASSPNP | (deferred) |
| [...] | | | |

- They import from each other extensively



Device driver image use-after-free

- In other words, drivers are resources that reference each other
 - refcounts!
- Each described by `LDR_DATA_TABLE_ENTRY`

```
kd> dt _LDR_DATA_TABLE_ENTRY
```

```
nt!_LDR_DATA_TABLE_ENTRY
```

```
[...]
```

```
+0x024 FullDllName      : _UNICODE_STRING
```

```
+0x02c BaseDllName     : _UNICODE_STRING
```

```
+0x034 Flags           : Uint4B
```

```
+0x038 LoadCount       : Uint2B
```

```
+0x03a TlsIndex        : Uint2B
```

16-bit only!



Device driver image use-after-free

```
C:\Users\test\Desktop>driverquery.exe
Name:      ntkrnlpa.exe, LoadCount: 110
Name:      halmacpi.dll, LoadCount: 89
Name:      kdcom.dll, LoadCount: 3
Name:      mcupdate_GenuineIntel.dll, LoadCount: 1
Name:      PSHED.dll, LoadCount: 3
Name:      BOOTVID.dll, LoadCount: 1
Name:      CLFS.SYS, LoadCount: 3
Name:      CI.dll, LoadCount: 2
Name:      Wdf01000.sys, LoadCount: 1
Name:      WDFLDR.SYS, LoadCount: 11
Name:      ACPI.sys, LoadCount: 1
Name:      WMILIB.SYS, LoadCount: 24
```

- If we load a driver that imports from e.g. `fwpkcInt.sys` 65,536 times, *LoadCount* is overflown.
 - must be a different path every time.
- Smallest default drivers take up 8kB - 65kB of virtual address space.
 - still within reasonable limits on X86-64 (within 4GB)

Device driver image use-after-free

- Not all drivers can be unloaded, even for `refcount=0`
 - there's a concept of `kernel DLLs`
 - not stand-alone, only loaded as dependencies
 - can be recognized by `DllInitialize` / `DllUnload` exports
 - examples: `usbport.sys`, `msrpc.sys`, `Classpnps.sys`
- Exploitation plan:
 - Find a small driver importing from a kernel DLL to load multiple times
 - Find another such driver which fails to load.
 - Overflow DLL refcount using driver A, then free using driver B.


Device driver image use-after-free

- Exemplary setting: use `wfplwf.sys` to overflow the `netio.sys` (DLL) refcount
- Use `tcpip.sys` to trigger the `free(netio.sys)`
- Works good!

Refcounts in the middle of an attack:

| | | | |
|--------|---------------|-------------|------|
| Name : | pcw.sys, | LoadCount : | 1 |
| Name : | Fs_Rec.sys, | LoadCount : | 1 |
| Name : | ndis.sys, | LoadCount : | 5656 |
| Name : | NETIO.SYS, | LoadCount : | 5655 |
| Name : | ksecpkg.sys, | LoadCount : | 1 |
| Name : | tcpip.sys, | LoadCount : | 1 |
| Name : | fwpmclnt.sys, | LoadCount : | 5640 |
| Name : | vmstorfl.sys, | LoadCount : | 1 |
| Name : | volsnap.sys, | LoadCount : | 1 |
| Name : | spldr.sys, | LoadCount : | 1 |

DLL modules
imported by
`wfplwf.sys`



Device driver image use-after-free

Effective result

[...]

<Unloaded_NETIO.SYS>+0x1b70:

88557b70 ?? ???

Resetting default scope

[...]

0: kd> kb

ChildEBP RetAddr Args to Child

8078a654 [...] <Unloaded_NETIO.SYS>+0x1b70

8078a668 [...] tcpip!CheckInboundBypass+0x1f

8078a810 [...] tcpip!WfpAleFastUdpInspection+0x55

[...]

Device driver image use-after-free

- Impact
 - Administrative rights required
 - Therefore, only `admin` → `ring-0` privilege escalation
 - Useful for subverting *Driver Signature Enforcement*
 - not much else

Device driver image use-after-free

Metrics

- Memory
 - `wfplwf.sys` takes 0x7000 bytes (28kB) of virtual memory.
 - 0x10000 (65,536) instances = ~2GB total.
- CPU time
 - Platform: Windows 7 64-bit, 4-core VMware Player, Intel i7-3930K @ 3.20GHz
 - ~100 loads per second.
 - 65,536 loads ~ 655 seconds ~ 10 minutes

win32k!NtUserCheckAccessForIntegrityLevel use-after-free

- On Wed, 30 Jun 2010 *Microsoft-Spurned Researcher Collective* dropped a 0-day at full disclosure.
 - Windows Vista / 2008 only
 - included a link to j00ru.vexillium.org :-/
- Turned out to be a trivial double-deref when accessing a *PsProcessType* object
 - Managed by the NT Object Manager

win32k!NtUserCheckAccessForIntegrityLevel

use-after-free

Faulty call chain

- win32k!NtUserCheckAccessForIntegrityLevel
 - win32k!LockProcessByClientId
 - win32k!LockProcessByClientIdEx
 - nt!PsLookupProcessByProcessId
 - nt!ObReferenceObjectSafe
 - nt!PsGetProcessSessionId
 - nt!ObfDereferenceObject
 - nt!ObfDereferenceObject

win32k!NtUserCheckAccessForIntegrityLevel

use-after-free

- Referenced once

`nt!PsLookupProcessByProcessId`

```
PAGE:006167A9  call     @ObReferenceObjectSafe@4
```

- Dereferenced twice

`win32k!LockProcessByClientId`

```
.text:BF88E63B  call     ds:__imp_@ObfDereferenceObject@4
```

`win32k!NtUserCheckAccessForIntegrityLevel`

```
.text:BF92D329  call     ds:__imp_@ObfDereferenceObject@4
```

- Broke the `reference# = dereference#` rule

win32k!NtUserCheckAccessForIntegrityLevel

use-after-free

- Bug allows arbitrary decrementation of `PointerCount` of an object.
- Conditions
 - Must be a process (`PsProcessType`)
 - In a different terminal session than caller (`process session id != gSessionId`)
 - System, `smss.exe`, `lsass.exe`, ...
 - Remote Desktop Services applications

win32k!NtUserCheckAccessForIntegrityLevel

use-after-free

- Exploitation concept
 - a. Find a process with `HandleCount = 0`
 - b. Free the object by dropping `PointerCount` to 0
 - c. Spray object memory with controlled data.
 - d. ???
 - e. PROFIT!
- `smss.exe` looks good

```
PROCESS 8c10b628  SessionId: none  Cid: 0194  Peb: 7ffda000  ParentCid: 0004
DirBase: 0015c020  ObjectTable: 87fc6fc8  HandleCount: 28.
Image: smss.exe
```

```
kd> !object 8c10b628
Object: 8c10b628  Type: (8465aec0) Process
ObjectHeader: 8c10b610 (old version)
HandleCount: 0  PointerCount: 22
```


win32k!NtUserCheckAccessForIntegrityLevel

use-after-free

Crash easy to trigger

```
TRAP_FRAME:  90706b0c -- (.trap 0xffffffff90706b0c)
ErrCode = 00000002
eax=86399708 ebx=8180c584 ecx=8c1232d0 edx=8c123310 esi=00000000 edi=00000000
eip=8187ec58 esp=90706b80 ebp=90706b88 iopl=0         nv up ei pl nz na po nc
cs=0008  ss=0010  ds=0023  es=0023  fs=0030  gs=0000             efl=00010202
nt!KiReadyThread+0x3c:
8187ec58 8906          mov     dword ptr [esi],eax  ds:0023:00000000=????????
Resetting default scope
STACK_TEXT:
[...]
90706b88 8188080e 819cfc20 863998ac 863998b4 nt!KiReadyThread+0x3c
90706ba4 818808d2 00000001 00000000 00000000 nt!KiUnwaitThread+0x14a
90706bc0 8187a307 00000001 8c1d0d78 863998ac nt!KiWaitTest+0xb6
90706bd8 81882cff 863998ac 00000001 00000001 nt!KeReleaseSemaphore+0x4f
90706c04 81d8d741 8c1d0f8c 00000001 00000000 nt!AlpcpSignalAndWait+0x7f
90706c40 81db91dc 00000001 90706cac 00000000 nt!AlpcpReceiveSynchronousReply+0x33
90706cd0 81dc041c 8c172818 00020000 00ddfab0 nt!AlpcpProcessSynchronousRequest+0x648
[...]
```

win32k!NtUserCheckAccessForIntegrityLevel use-after-free

- Exploitation more difficult
 - Only candidate is `smss.exe` (despite *System*)
 - Unknown `PointerCount`
 - Requires advanced kernel pool feng-shui
 - `EPROCESS` takes 0x25c (604) bytes of `NonPagedPool`
 - failed attempt = Blue Screen of Death
- Definitely still possible!
 - keep an eye on my blog ☺

win32k!NtUserCheckAccessForIntegrityLevel

use-after-free

- Impact
 - Local privilege escalation if exploitation succeeds
 - Denial of Service otherwise.
 - Windows Vista / 2008 Server only.
- Metrics
 - Memory: irrelevant
 - CPU time: irrelevant (instant)
- Fix
 - Setting the output object pointer to NULL in win32k!LockProcessByClientId
 - second dereference doesn't occur anymore


win32k!NtUserAttachThreadInput

use-after-free

- Some threads in Windows are marked as GUI
 - can then talk to `win32k.sys`
 - required for anything graphics-related
- Every such thread has a kernel-mode message queue.

```
kd> dt tagQ
win32k!tagQ
    +0x000 mInput          : tagMLIST
    [...]
    +0x13c QF_flags        : Uint4B
    +0x140 cThreads        : Uint2B
    +0x142 cLockCount      : Uint2B
    +0x144 msgJournal      : Uint4B
```

looks like
refcounts!



win32k!NtUserAttachThreadInput

use-after-free

- Threads can attach to each others' queues!
 - see [AttachThreadInput](#) (documented API)
- Queues must store # of reliant threads
 - uses cThreads for just that
- Queues freed in win32k!UserDeleteW32Thread when (cThreads == 0) && (cLockCount == 0)

```
.text:BF8D6B63  cmp     [ecx+tagQ.cLockCount], di
.text:BF8D6B6A  jnz     short loc_BF8D6B7D
.text:BF8D6B6C  mov     eax, ecx
.text:BF8D6B6E  cmp     [eax+tagQ.cThreads], di
.text:BF8D6B75  jnz     short loc_BF8D6B7D
.text:BF8D6B77  push    eax                ; Entry
.text:BF8D6B78  call    _FreeQueue@4       ; FreeQueue(x)
```

win32k!NtUserAttachThreadInput

use-after-free

- There's no refcount leak in the implementation
 - no "free" incrementations
- Can we legitimately attach > 65,535 threads to a single queue?
 - Yes, if we can create that much.
 - Can we?
- Mark Russinovich had an excellent post about it, see *"Pushing the Limits of Windows: Processes and Threads"*

win32k!NtUserAttachThreadInput

use-after-free

- Short answer: **no** on 32-bit Windows
 - limitations: kernel virtual address space size, physical memory capacity, ...
 - only up to 32K threads, usually far less.
- Good news: **yes** on 64-bit Windows

win32k!NtUserAttachThreadInput

use-after-free

Let's test!

```
for (unsigned int i = 0; ; i++) {  
    if (!CreateThread(NULL, 0,  
                      (LPTHREAD_START_ROUTINE) ThreadRoutine,  
                      NULL, 0, NULL)) {  
        break;  
    }  
    printf(„threads: %u\n“, i);  
}
```


win32k!NtUserAttachThreadInput

use-after-free

```
c:\code\testlimit\objchk_win7_amd64\amd64>test  
threads: 157179  
c:\code\testlimit\objchk_win7_amd64\amd64>
```

- Windows 7 64-bit, 12GB of RAM
- ~ 2.64 GB physical memory consumption for 65,536 threads
- Several seconds of CPU time

win32k!NtUserAttachThreadInput

use-after-free

Security by poor programming practices?

- Overflowing a 16-bit counter shouldn't take too long, right?
 - in theory...
- Every "*attach thread A to B*" request:
 - results in a full recalc of thread queues
 - takes $O(n^2)$ time, n = session thread count
- Creating a queue with 2^{16} threads takes $\sim 2^{48}$ steps
 - could've been done a whole lot faster

win32k!NtUserAttachThreadInput

use-after-free

AttachThreadInput (x, y) algorithm (pseudo-code)

```
win32k!gpai.append(pair(thread_from, thread_to));
foreach thread in current_thread->desktop:
    pqAttach = thread->pq;
    changed = false;
    if thread->attached:
        continue
    do:
        foreach thread_nested in current_thread->desktop:
            if thread_nested->pq == pqAttach:
                foreach req in win32k!gpai:
                    if req.first == thread_nested || req.second == thread_nested:
                        attach(req.first, req.second)
                        changed = true
while changed;
```

win32k!NtUserAttachThreadInput

use-after-free

- Still exploitable (with some extra work)
 - **Note:** recalc only for caller thread's desktop
- Plan:
 - Create `self_desktop` and `thread_desktop` desktops
 - Assign main thread to `self_desktop`
 - Create 65,536 threads
 - assign all to `thread_desktop`
 - Attach threads 1..65,536 to 0
 - fills in the `win32k!gpai` list with thread pairs
 - fast: single attach is $O(1)$ for foreign desktops (no recalc)

[...]

win32k!NtUserAttachThreadInput

use-after-free

- Plan, part two

- Switch main thread and current workstation to `thread_desktop`
- Attach main thread queue to thread 0
 - causes a full recalc, $n = 2^{16}$, $O(n^2) \sim 2^{32}$ iterations
 - within one syscall, no context switches
 - triggers the integer overflow; `refcount = [...], 65536, 0`
 - triggers a free of the shared input queue
- Spray session paged pools
- Terminate remaining threads
 - triggers use of the freed queue

win32k!NtUserAttachThreadInput

use-after-free

Results

- Multiple assertion hits on a checked build

```
(s: 1 0x4dc.484 test.exe) [Err] DBGValidateQueueStates: Assertion failed: (pti == pq->ptiMouse)
|| (fAttached && (pq == pq->ptiMouse->pq))

(s: 1 0x4dc.484 test.exe) [Err] DBGValidateQueueStates: Assertion failed: (pti == pq-
>ptiKeyboard) || (fAttached && (pq == pq->ptiKeyboard->pq))
```

- Ultimately, a bugcheck

```
win32k!DestroyThreadsMessages+0x22:
fffff960`0011a6b6 488b33 mov rsi,qword ptr [rbx] ds:002b:aaaaaaaa`aaaaaaaa=????????????????
Resetting default scope

STACK_TEXT:
fffff880`fd18e7d0 fffff960`00119da9 : [...] : win32k!DestroyThreadsMessages+0x22
fffff880`fd18e800 fffff960`0013deb7 : [...] : win32k!xxxDestroyThreadInfo+0x1001
fffff880`fd18e8d0 fffff960`00115140 : [...] : win32k!UserThreadCallout+0x93
fffff880`fd18e900 fffff800`0299d375 : [...] : win32k!W32pThreadCallout+0x78
```

win32k!NtUserAttachThreadInput

use-after-free

- Impact
 - Invincible processes by *infinite* loops in win32k.sys
 - Denial of Service (failed use-after-free exploitation)
 - Escalation of Privileges (successful exploitation)
 - resource constraints
 - kernel pool feng-shui required again
- Metrics
 - Memory: ~2.5GB required for thread storage.
 - CPU time: up to 10 minutes
 - creating threads (2^{16} steps): < 5s
 - attaching threads (2^{16} steps) < 2 minutes
 - doing global recalc (2^{32} steps) < 10 minutes

win32k!NtUserAttachThreadInput

use-after-free

The fix

- Expand the `cThreads` / `cLockCount` refcounts to 32 bits
 - you can't possibly have 4,294,967,296 threads... yet (but ping me when you can)

```
[...]
+0x140 cThreads      : Uint2B
+0x142 cLockCount    : Uint2B
[...]
```



```
[...]
+0x140 cThreads      : Uint4B
+0x144 cLockCount    : Uint4B
[...]
```

```
[...]
add      word ptr [r12+140h], 1
[...]
```



```
[...]
inc      dword ptr [r12+140h]
[...]
```


win32k!NtGdiAddFontResource use-after-free

- Applications can load external fonts for local usage
 - documented AddFontResource Windows API
 - perhaps used in every win32k.sys font fuzzer
- " When an application no longer needs a font resource it loaded by calling the AddFontResourceEx function, it must remove the resource by calling the RemoveFontResourceEx function.
- "
- Sounds reference-countable! :-)

win32k!NtGdiAddFontResource use-after-free

- Indeed...

Callstack


```
kd> kb
ChildEBP RetAddr  Args to Child
9b714af4 [...] win32k!PFFOBJ::vLoadIncr+0x12
9b714b14 [...] win32k!PFTOBJ::chpfeIncrPFF+0x94
9b714b80 [...] win32k!PUBLIC_PFTOBJ::bLoadFonts+0x90
9b714bc8 [...] win32k!GreAddFontResourceWInternal+0xad
9b714d14 [...] win32k!NtGdiAddFontResourceW+0x15e
9b714d14 [...] nt!KiFastCallEntry+0x12a
0022fd2c [...] ntdll!KiFastSystemCallRet
```

win32k!NtGdiAddFontResource use-after-free

- Indeed...

```
.text:BF8149BF ; public: void __thiscall PFFOBJ::vLoadIncr(unsigned long)
[...]
```

| | | |
|------------------------------|---------------------------------|-----------------------|
| .text:BF8149C4 | test | [ebp+arg_0], 20h |
| .text:BF8149C8 | mov | eax, [ecx] |
| .text:BF8149CA | jz | short loc_BF8149D1 |
| .text:BF8149CC | inc | dword ptr [eax+28h] |
| .text:BF8149CF | jmp | short loc_BF8149D4 |
| .text:BF8149D1 | | |
| .text:BF8149D1 loc_BF8149D1: | | |
| .text:BF8149D1 | inc | dword ptr [eax+24h] |
| .text:BF8149D4 | | |
| .text:BF8149D4 loc_BF8149D4: | | |
| .text:BF8149D4 | call | PFFOBJ::vRevive(void) |
| .text:BF8149D9 | pop | ebp |
| .text:BF8149DA | retn | 4 |
| .text:BF8149DA | ?vLoadIncr@PFFOBJ@@QAEXK@Z endp | |



refcount
incrementation!

win32k!NtGdiAddFontResource use-after-free

- Details
 - 32-bit refcount involved on both X86 / X86-64
 - perhaps an ULONG, but exact structure unknown
 - No persistent memory allocations!
- How long does it take?
 - well, 2^{32} system calls...
 - test environment: Windows XP SP3 in a VM, single core
 - incr. rate at about 100,000 requests / second
 - (only) ~12 hours!
 - could be less on better machine or with optimized exploit

win32k!NtGdiAddFontResource use-after-free

- Results

- Upon unload, the `PFFOBJ` class is "killed" when refcount drops to 0.
- Stack trace:

```
#0 win32k!PFFOBJ::vKill
```

```
#1 win32k!PFFOBJ::bDeleteLoadRef
```

```
#2 win32k!PFTOBJ::bUnloadWorkhorse
```

```
#3 win32k!GreRemoveFontResourceW
```

```
#4 win32k!NtGdiRemoveFontResourceW
```

win32k!NtGdiAddFontResource use-after-free

- All sorts of badness
 - use-after-frees
 - NULL pointer dereferences

```
kd> g
Access violation - code c0000005 (!!! second chance !!!)
win32k!bGetNtoD_Win31+0x1f:
82008864 8b4830          mov     ecx,dword ptr [eax+30h]
kd> ? eax
Evaluate expression: 0 = 00000000
kd> kb
ChildEBP RetAddr  Args to Child
9bb28bc4 [...] win32k!bGetNtoD_Win31+0x1f
9bb28bf8 [...] win32k!PFEOBJ::bSetFontXform+0x3e
9bb28c98 [...] win32k!RFontOBJ::bInit+0x1bf
9bb28cb0 [...] win32k!RFontOBJ::vInit+0x16
9bb28cd4 [...] win32k!GreGetRealizationInfo+0x2a
9bb28d24 [...] win32k!NtGdiGetRealizationInfo+0x41
9bb28d24 [...] nt!KiFastCallEntry+0x12a
```

win32k!NtGdiAddFontResource use-after-free

- Impact
 - typically DoS or EoP, depending on exploitation skills
 - works on 32-bit and 64-bit platforms
- Fix: mount a reference count limit at ULONG_MAX
 - Quite risky, what if there's a two-thread race?

win32k!NtGdiAddFontResource use-after-free

```
v2 = *(_DWORD *)this;  
if ( a2 & 0x20 )  
    ++*(_DWORD *)(v2 + 40);  
else  
    ++*(_DWORD *)(v2 + 36);  
return PFFOBJ::vRevive();
```



```
v2 = *(_DWORD *)this;  
if ( a2 & 0x20 )  
    v3 = v2 + 40;  
else  
    v3 = v2 + 36;  
if ( *(_DWORD *)v3 == -1 )  
{  
    result = 0;  
}  
else  
{  
    ++*(_DWORD *)v3;  
    PFFOBJ::vRevive();  
    result = 1;  
}
```

overflow
prevention
check

hmmm... a
new bug?

Also worth checking out

CVE-2011-2013 (tcp/ip stack use-after-free)

- Fixed on November 8, 2011
- 32-bit reference counter integer overflow
- Remote, through UDP packets!
- Works on closed ports!
- Root cause - adverse circumstances and no mitigations
 - "small" integer, a 64-bit one would suffice
 - no sanity checks
 - no persistent memory allocations bound to refcount incrementations

Mitigations concepts

Preventing refcount problems

- You can't prevent developers from writing buggy code
- But you can mitigate consequences of the resulting vulns
 - Provide a "secure" interface for everyone to use
 - Not perfect, but raises the bar

Preventing integer overflows

- Introduce `refcount_t` as an alias to `int64_t`
 - doesn't cost anything: memory is cheap
 - times when it mattered are long gone
 - would prevent 99% refcount overflow attacks
 - **potential problem**: sometimes counters are difficult to recognize

Preventing integer overflows

- Introduce generic APIs for refcount manipulation
 - `nt!IncrementRefcount`, `nt!DecrementRefcount`, `nt!TestRefcount`
 - could include basic sanity checks

```
if (++(*refcount) < 1) {  
    KeBugCheckEx(REFCOUNT_GONE_WRONG);  
}
```

```
if (--(*refcount) < 0) {  
    KeBugCheckEx(REFCOUNT_GONE_WRONG);  
}
```

Preventing excessive dereferences

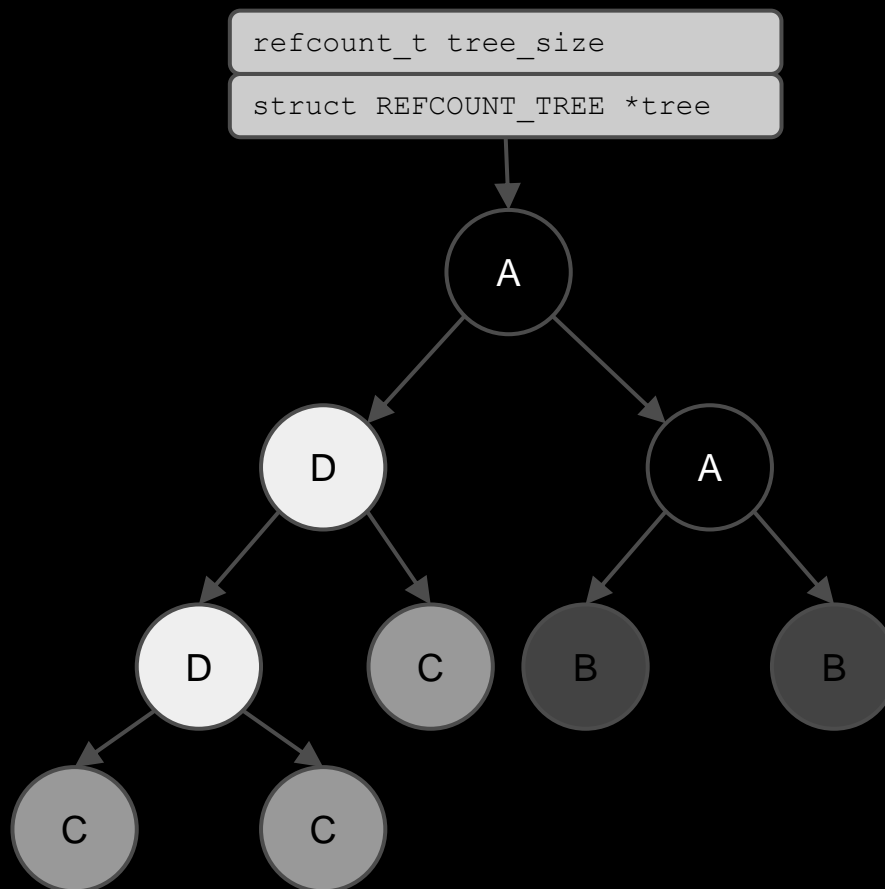
- Much harder than plain integer problems
 - let's never free refcounted allocations! :-)
 - revisit the idea when unlimited memory available
 - curio: *nt!NtMakePermanentObject*
 - requires *SeCreatePermanentPrivilege*
- The interface doesn't know caller's logic
 - which derefs are paired with which refs?

Preventing excessive dereferences

- Idea: identify each ({reference}, {dereference}) pair with a unique tag
 - similarly to *pool tags*
- A "reference counter" becomes a "reference tree"
- Store information about all pending reference tags in the tree
- Always pass the tag to the {ref,deref} API
 - Test if tag is in tree before decrementing

Preventing excessive dereferences

- A self-balancing binary search tree



AVL trees already
implemented in Windows

Preventing excessive dereferences

Performance hit

| | Tree implementation | Traditional implementation |
|-------------------------|---------------------|----------------------------|
| <i>Reference cost</i> | $O(\lg n)$ | $O(1)$ |
| <i>Dereference cost</i> | $O(\lg n)$ | $O(1)$ |
| <i>Test cost</i> | $O(1)$ | $O(1)$ |

- Statistics (Windows 7 SP1 32-bit with a few apps)
 - Average *PointerCount*: $118009 / 29364 \approx 4.01883$
 - Average *HandleCount*: $15135 / 29364 \approx 0.51542$
- Difficult to measure refs/derefs per second
- Overhead should be acceptable (own opinion)

Preventing excessive dereferences

Memory overhead

- Loose estimate
 - ~120,000 references to NT executive objects at startup
 - Twice as much during typical session =~ 250,000
 - Twice as much including other refcounts =~ 500,000
 - Assume 64 bytes per one reference
 - pool header, tag, pointers to parent / children
 - A total of extra ~30MB of NonPaged memory
 - guess if my 12GB RAM machine can take it?

Preventing excessive dereferences

Other problems

- Lazy developers
 - would have to define unique tags
 - already do it for pool allocations, so perhaps possible?
- Legacy issues
 - existing API routines lack tagging information
 - `ObReferenceObject{ByHandle,ByPointer}`
 - how to communicate failure (e.g. lack of memory)?
- Passing tags through wrappers
- Possibly low engineering effort / benefit ratio
 - how many bugs would this prevent?

Preventing excessive dereferences

Benefits

- If properly executed, would prevent most use-after-frees through double derefs
 - *stealing* references not possible anymore
 - *dereference* sequence would have to match the *reference* one to exploit
- Automatic mitigation integer overflow
 - through memory constraining
- Robust interface for future use

Conclusions

Random thoughts

- Refcounts bugs = use-after-frees
 - otherwise rarely observed (perhaps except Tarjei)
 - usually time-consuming and tricky to exploit
 - often memory-consuming
- Kernel pool spraying should be better investigated
- Integer types \neq machine word don't scale
 - No explicit (1/2 void*) or (1/4 void*)
 - Small types used 20 years ago can take revenge
 - More to be found?

Random thoughts

- Inconsistent patches
 - sometimes extending types
 - sometimes pinning
 - sometimes sanity checks
 - would a common interface help?
- Microsoft doesn't backport fixes?
 - Why [CVE-2010-2549](#) only affected Vista / 2008?
 - Could've been found by bindiffing?
 - See Nikita's talk

Благодарю вас за
внимание!

Questions?



E-mail: j00ru.vx@gmail.com

Blog: <http://j00ru.vexillium.org/>

Twitter: [@j00ru](https://twitter.com/j00ru)